



THE FIELD GUIDE

Know Your AgenticAi

The 5 questions to answer before you let an AI agent act on your behalf.

AI stopped waiting for instructions. It started acting. This guide is the fastest way to find out whether the agents in your business are provable, auditable, and accountable, before a regulator, a client, or a court asks you to prove it.

Patent-pending framework · UK IPO · GB2611754.9

Rume Dominic · AI Engineer, Founder of VOREM · 7+ years shipping autonomous agents
rumedominic.com

The question nobody is asking

A traditional model gives an answer and stops. An agent takes an action.

It sends the email. It moves the money. It approves the claim. It changes the record. It does this across your tools, your data, and your customers, on its own. That is agentic AI, and it is already inside most businesses that touch software.

Everyone is asking the same two questions: how smart is it, and how fast can we deploy it. Almost nobody is asking the one that actually decides whether you are safe: **can you prove what it did, and why?**

When an agent makes a mistake, and it will, most teams cannot reconstruct what happened or name who is accountable. That is not a technology problem. It is a liability sitting on your balance sheet, invisible until someone asks you to explain a decision your AI made and you cannot.

In banking we solved a version of this decades ago. Before you move money for a stranger, you have to Know Your Customer. Identity. Provenance. A record you can show a regulator. Agentic AI needs the same discipline. I call it Know Your AgenticAI.

01

Provable

You can show, after the fact, exactly what the agent did and why.

02

Auditable

An outsider can inspect that record and trust it.

03

Accountable

When it goes wrong, there is a clear line to a named human.

An AI you cannot audit is dangerous no matter how smart it is, and no matter who built it. The five questions on the next page turn that principle into a test you can run today.

The 5 questions

Sit your team down and answer these out loud, in writing, for one agent you have in production. If you cannot answer them, you are not ready to let it act.

QUESTION 1

What is this agent authorised to decide alone?

Be specific. "It handles support" is not an answer. "It can issue refunds under \$50, and nothing else" is. An agent without a clear boundary will find the edge of it at the worst possible time.

QUESTION 2

What requires human sign-off?

Every high-stakes action needs a gate. Money above a threshold. Anything legal. Anything that touches a customer's data or reputation. Decide this before deployment, not after an incident.

QUESTION 3

If it makes a mistake, how do you find out?

Not "if". When. Do you log the agent's reasoning, not just its output? If the honest answer is "we would notice eventually", you do not have an answer.

QUESTION 4

If it makes a mistake, who is accountable?

An agent cannot be fired, sued, or held responsible. A human owns every action it takes. Name that person now, while it is calm, not during a crisis.

QUESTION 5

Can you prove to a regulator what it did and why?

This is the one that ends careers. If a regulator, a client, or a court asks you to explain a decision your AI made, can you produce a clear, trustworthy record? If not, you are carrying a risk you cannot see.

The 60-second assessment

For each question, can you answer it clearly and back it with a record? Tick one box per row.

Question	Yes, with a record / No
1 Authorised decisions are defined	☐ ☐
2 Human sign-off gates exist	☐ ☐
3 Agent reasoning is logged	☐ ☐
4 A named human is accountable	☐ ☐
5 You could prove it to a regulator	☐ ☐

Read your score.

Five yeses: rare. You are ahead of almost everyone deploying agents today.

One or two nos: you have exposure you can close quickly. Start with the reasoning log.

Three or more nos: you are running an agent you cannot prove, audit, or account for. That is the risk this framework exists to remove, before it finds you first.

Most companies deploying agents today would fail three of the five. They optimise for speed and treat governance as paperwork. It is the opposite. Governance is what lets you move fast without blowing up.

TAKE IT FURTHER

Turn the checklist into certainty

This guide is the free version of a discipline I apply for a living. The KYA Audit is a fixed-scope assessment of one of your agents against the full standard: a provability report, an auditability score, a named accountability chain, a regulator-readiness verdict, and a 90-day roadmap to close every gap.

If the audit does not surface at least three material gaps you were not already tracking, it is free. You keep the report either way.

[Book a call at rumedominic.com](https://rumedominic.com)

Let us build AI we can prove.

Rume