

ORume Dominic Urie (Rume Dominic)

Daml and Canton Engineer

Birmingham, UK · orumedominic@gmail.com · +44 7402 830944 · linkedin.com/in/dominicrume · github.com/dominicrume · rumedominic.com

Eight years in blockchain. I build the part of a system that proves what it **refused** to do. On Polygon I anchored what AI agents did. On Canton, in Daml, I also record what they were stopped from doing, with a test suite that breaks its own checks to show they hold.

CANTON AND DAML

- **3rd place, Cantor8 “Build on Canton” hackathon**, London, August 2026. Team Vorem, \$1,500. Judged by a panel from across the Canton ecosystem on technical, security, institutional and commercial grounds.
- **Know Your AgenticAI**. Records what an AI agent was **refused** as a contract on the ledger instead of a log line in the operator’s own system. Daml package **kya-rails-mandate**, SDK 3.4.11, vetted on Canton DevNet. 111 Daml test scripts, 32 spending fences, and 82 mutations that each break something real and require a named test to go red.
- **3 pull requests merged into the Canton Developer Hub** (#156, #160, #168). One fixed link rendering for nineteen catalogue entries, eighteen of which belonged to other people. Listed in the official Canton catalogue.
- **4 pull requests filed into the Cantor8 hackathon toolkit** (#7 to #10). Access token cached past expiry; **find_party** scanning every party on the validator; **C8_USER** read at import rather than at call time; a stale README.
- **3 findings reports to OpenZeppelin’s Canton libraries**. 52 of 69 authorisation fences could be made vacuous with every test still passing.
- **Dev Fund proposal** open under RFP 27, security monitoring and auditability (**canton-dev-fund #860**).

PUBLISHED SOFTWARE

- **knowyouragenticai-receipts** (PyPI). A receipt format with a written specification and 20 conformance vectors. Three independent implementations, in Python, JavaScript and Go, agree on all 20. MIT, zero dependencies, 241 downloads a month.
- **ai-code-quality-auditor** (PyPI). Measures AI coding agents for hallucination, security and code quality. The core instrument of my MSc research. 472 downloads a month.

SELECTED ENGINEERING

- **AgenticOS, Veritaport Ltd** (veritaport.co.uk). The evidentiary layer of a forensic system, in Solidity on Polygon. Server-side signed record stamping with on-chain anchoring, so a third party verifies record integrity in seconds without per-user wallets.
- **Matty’s Place tenant hub, via MCKI Solutions** (app.mattysplace.org.uk). Enterprise HMO platform. 100+ tenants, multi-brand, role-based access architecture, SHA-256 audit chain, and every AI agent action anchored to Polygon. Saved roughly 100 hours a week and £80,000 a year.
- **Quant trading analytics platform**. Five-model consensus engine with live subscribers and revenue.
- **Nigeria tax calculator**. Retrieval and OCR system reading bank statements and classifying transactions for tax. Deployed.

EDUCATION

- **MSc Artificial Intelligence and Business Strategy**, Aston University, 2026. Dissertation: measuring the quality and safety of AI-generated code.

ALSO

- Pending UK patent in verifiable AI accountability, **GB2611754.9**. 12,000+ people trained across four continents in eight years. Five published books. 28+ public repositories.

Available now, full time or contract, UK based. Every number above is public and checkable.